

Yaoshuo Wu 吴耀铄

✉ yaoshuowu75@gmail.com 🌐 potato0705.github.io

🌐 github.com/Potato0705 📍 Shenzhen, China



RESEARCH INTERESTS

- **AI in educational assessment:** validity and reliability of LLM-based scoring, including robustness to gaming
- **Educational agents:** teaching agents evaluated on real learning outcomes rather than predictive accuracy
- **Labor-market data for curriculum planning:** identifying genuinely emerging skill requirements

EDUCATION

Shenzhen Technology University (SZTU)

Shenzhen, China

B.Eng. in Intelligent Perception Engineering, College of Engineering Physics

2023 – 2027 (expected)

GPA 3.78/4.5 · Average score 85.93/100 · Major rank 11/138

Honors: Provincial First Prize, Lanqiao Cup (16th and 17th editions)

2025, 2026

PUBLICATIONS

Accepted

[†] corresponding author

- [1] **Yaoshuo Wu**, Haiqin Yang[†], Zhong Ming[†]. Opinion as Prior: Supervised VA Prediction for DimASQP. *ICONIP 2026 (Springer CCIS)*, accepted. [Code]

Under Review

- [2] **Yaoshuo Wu** et al. Accuracy Is Not Validity: Construct Validity Auditing for LLM Essay Graders. *Under review at AAAI-27*.
- [3] **Yaoshuo Wu**, Haiqin Yang[†], Yue Wang, Irwin King. Breadth over Trend: Validating Emerging Job Requirement Signals with the Bilingual EJRD Benchmark. *Under review at the KDD 2027 Datasets and Benchmarks Track*. [Benchmark & Code]
- [4] **Yaoshuo Wu**, Shuxian Hou, Xiaochun Chen, Danyun Chen, Xiaomin Li, Jiahui Zhou, Shenyu Liu, Haiqin Yang[†]. Low-Resource Multilingual Event Mining: A Survey and a Framework for Evaluation and Reporting. *Under review at ACM Computing Surveys*. [MELEE Toolkit]

RESEARCH EXPERIENCE

Undergraduate Research, Shenzhen Technology University

Dec 2025 – Present

Focus: NLP and the validity of AI evaluation, recently applied to education

- Proposed the Construct Validity Margin (CVM) to audit LLM essay graders; on ASAP, 18 of the 27 grader-prompt cells with $QWK \geq 0.5$ respond at least as much to gaming edits as to content loss [2].
- Introduced EJRD, a bilingual benchmark of 1,725 job requirements from 221,582 Hong Kong postings with a 500-item human gold set, and SAVA, a framework that validates signals only against gold standards they did not define [3].
- Proposed a construct-validity-based evaluation and reporting framework for low-resource multilingual event mining, backed by a double-coded audit of 39 studies; open-sourced the MELEE toolkit [4].
- Proposed opinion-prior supervision for valence-arousal (VA) prediction in dimensional aspect-based sentiment analysis, lowering mean Euclidean VA error by 3.4% (Restaurant) and 7.5% (Laptop) relative to a token-level baseline [1].

COMPETITIONS

CyberAI Cup 2026, Team SZTU-AGI

Jul – Aug 2026

- **Task 1 (packaging material difference mining):** to localize small content changes between a package template and its photograph amid image-wide irrelevant differences, developed a pipeline that registers the photograph to the template, feeds signed high-pass difference maps to a U-Net, and fuses a three-seed ensemble with a teammate's TASL model.
- **Task 2 (farm compliance audit):** to judge 41 checkpoints per farm without labeled data, where “not applicable” requires positive evidence, helped build the team's LLM pipeline with policy-then-evidence retrieval, cited verdicts for each regulatory group, and a separate applicability gate.
- **Task 3 (encrypted RTC traffic identification):** to identify a flow's application and call mode (voice or video) from only the sizes and timings of its first five packets, developed a pipeline with 51 engineered features, a LightGBM/XGBoost/TabICL v2 ensemble, and class-prior correction for class imbalance.
- Wrote the team's final technical report.